

EVIDENCEFIRST: Evidence-State Monitoring for Auditable GraphRAG Systems

Anonymous Author(s)

Institution anonymized for review
email@anonymized.com

Abstract. Fixed-context Web GraphRAG systems combine Web pages, graphs, and LLM readers, but operators still need to know which evidence condition preceded a suspicious answer. EVIDENCEFIRST is an evidence-state monitoring and auditing framework for GraphRAG systems. Before generation, a deterministic checker assigns each question to a checked candidate path, repaired path, residual graph gap, or passage-fallback state; the saved state records the repair or fallback route, reader evidence, selector decision, and audit fields. The objective is evidence-state analysis rather than leaderboard optimization. Across HotpotQA-1000 and 2WikiMultiHopQA-500, EVIDENCEFIRST remains within the operating range of recent GraphRAG baselines while exposing audit-risk queues that can be replayed without regeneration. Empirically, saved audit signals form review queues whose top-20% strata reach 67.0% and 40.0% error rates on HotpotQA and 2Wiki; adding graph-state fields to answer-selection signals improves error-detection AUC (0.6077→0.6321 on HotpotQA) and, more importantly, supplies typed routing causes: inside selection-conditioned queues, selected unrepaired-incomplete Hotpot cases and non-selected residual-gap 2Wiki cases are 26.6 and 33.3 error points worse than their respective comparison strata. Rather than predicting why a GraphRAG system failed, EVIDENCEFIRST records a replayable evidence-state trace that helps failures be routed, audited, and recomputed after deployment.

Keywords: Knowledge Graphs · RAG on the Web · GraphRAG · Evidence Observability · Auditing · Routing

1 Introduction




Fig. 1. Evidence-state pipeline and state-machine transitions. The Evidence-State Checker runs a structural state machine whose terminal state is saved per-question and written to the audit trace for replayable inspection.

GraphRAG research has mainly optimized retrieval, reasoning, and generation over retrieved text and graph structure. In an enterprise search or public Web QA service, however, the operational question after a bad answer is often not “which model scored higher?” but “which evidence condition preceded this case?” For example, a Web QA service answering from search-result pages must, after a complaint, determine whether the retrieved pages lacked a bridge entity, the constructed graph was disconnected, or the reader misread connected evidence – distinctions invisible in answer-plus-context logs. The failure may come from missing retrieved context, an incomplete graph, an absent bridge entity, disconnected evidence, a reader mistake, or final answer selection. RAG grounds language-model answers in retrieved text [13], while KG-RAG and GraphRAG systems add entity-relation structure over retrieved corpora [5, 9, 11]. Yet most systems expose only the final answer and retrieved context, leaving operators without a stable evidence-state trace for routing and audit.

EVIDENCEFIRST treats the evidence graph as an observable state object before the reader is called. The checker assigns each question to a checked candidate path, repaired path, residual graph gap, or passage-fallback state. That state is operational rather than merely logged: it selects graph repair or passage fallback, constrains which evidence is passed to generation, records the downstream selection decision, and can be replayed from saved artifacts.

The objective of this work is evidence-state analysis rather than leaderboard optimization. We therefore evaluate on representative subsets that allow detailed replay inspection and audit analysis. HotpotQA and 2Wiki provide retrieved passages, entity mentions, and linked evidence graphs; EVIDENCEFIRST studies how a GraphRAG system exposes and uses the retrieved graph state before generation.

Contributions. This paper makes three contributions:

- **Evidence-state abstraction.** We introduce a lightweight abstraction that ties each answer artifact to a checked, repaired, residual, or fallback graph state computed before generation.
- **Replayable audit protocol.** We support recomputation of reported claims, risk strata, path-state summaries, and evidence-visibility diagnostics from saved artifacts without regeneration.
- **Failure routing mechanism.** We study how graph-state signals facilitate repair and inspection workflows, separating graph repair, reader/selector inspection, and context review.

2 Related Work

RAG, GraphRAG, and multi-hop retrieval. RAG conditions a language model on retrieved evidence [13], while GraphRAG systems organize evidence as entity, relation, or community structures [5, 9, 11]. Hierarchical, memory-oriented, and reasoning-aware variants structure context as trees, long-term memories, hybrid records, or iterative retrieval paths [21, 10, 15, 23, 4, 16]. These systems mainly acquire, traverse, or compress evidence for a reader. EVIDENCEFIRST targets a later operational stage: after graph evidence is retrieved or constructed, it records whether the graph is checked, repaired, residual, or fallback before generation.

Observability and traceability in information systems. Information systems expose internal state for later inspection: databases preserve transaction logs [17], distributed systems expose execution traces [22], and data-management systems track provenance or lineage [1, 3]. GraphRAG systems have an analogous need, but saved artifacts usually stop at retrieved evidence, generated text, or tool-call logs. EVIDENCEFIRST imports this observability perspective into Web GraphRAG by recording the graph condition that routed each answer artifact.

Verification and auditability. RAG evaluation frameworks score retrieval, generation, and answer quality [6, 20, 2, 7], while attribution studies ask whether claims are grounded in identifiable sources [18, 8, 24]. S2G-RAG emits text gap items [14], and RAGChecker defines diagnostic metrics for retrieval and generation modules [19]. EVIDENCEFIRST narrows the observable object to a per-question graph state stored before generation, used for repair/fallback routing, and replayed from saved artifacts.

3 Method

3.1 Web Knowledge-RAG Problem Setting

We model a Web GraphRAG query as fixed-context inference over retrieved or provided Web evidence. Given a question q , candidate passages $\mathcal{P} = \{p_i\}$, and a per-question evidence graph $\mathcal{G} = (V, E)$ built from the same visible evidence, the system outputs a short answer a and a diagnostic record D . Nodes in $\mathcal{G}$

are entities or evidence units, and edges are extracted relations, support links, or passage-level associations. The setting covers search results, Wikipedia-style pages, entity descriptions, and linked Web documents; it does not assume live crawling or changing Web content during inference.

EVIDENCEFIRST records a structural path-state contract before generation. The graph state is one of four states: a checked candidate path, a repaired path after localized graph expansion, a residual graph gap after unsuccessful repair, or a passage-fallback state when usable graph evidence is absent. The diagnostic record D stores path-checked status, path length, gap type, missing-entity and disconnected-pair counts, repair/fallback tags, and the final answer-selection flag. The recorded state authorizes the repair, fallback, or evidence route that generation receives.

Table 1. Evidence-state contract used by EVIDENCEFIRST. The states are operational routing states over the retrieved graph.

State	Entry condition	Operational route	Saved fields
Checked path	Candidate path satisfies entity mapping, answer-type filter, and minimum path length	Terminal reader route: pass extracted chain, with local passages when enabled	<code>chain_complete</code> , chain length, empty gap label
Repaired path	Initial gap triggers entity-augmentation or bridge repair and recheck becomes complete	Transition from gap to terminal reader route; preserve repair provenance	repair tags, <code>chain_complete</code> , chain length
Residual gap	Recheck remains incomplete or short after repair	Terminal reader route with available triples/passages and retained graph cause	gap label, missing entities, disconnected-pair count, repair tags
Passage fallback	No usable triples or graph evidence is absent	Terminal passage route; mark graph evidence unavailable	fallback tags, empty/unknown chain fields

Algorithmic contract. Given q , $\mathcal{P}$, and $\mathcal{G}$, EVIDENCEFIRST executes a fixed state machine before the reader output. It extracts question entities, maps them to graph nodes by exact, substring, and token-overlap matching, selects candidate answer nodes, filters them by coarse answer type, and tests shortest paths in the evidence subgraph. If an adequate path exists, the state is checked and the extracted chain is sent to the reader. Otherwise, the checker emits `missing_entities`, `disconnected`, `bridge`, or `short_chain`; empty graph evidence becomes passage fallback. Missing entities trigger entity-augmentation repair, disconnected pairs trigger bridge repair over at most five endpoint pairs, and every repair is followed by the same structural recheck. A complete recheck becomes a repaired path; otherwise the residual or fallback state is retained through generation and selection. These rules are fixed before evaluation and use no gold answers or support facts at test time.

3.2 Evidence-First Inference

EVIDENCEFIRST performs inference by applying the same state contract to every question: it builds or loads a per-question KG, retrieves triples and passages, applies the pre-reader structural check, optionally repairs and rechecks localized gaps, and generates from the evidence packet selected by the resulting state. Empty graph evidence routes to passages instead of being silently treated as complete, and residual gaps remain attached to the final answer record.

3.3 Path-State Checking and Diagnostic Repair

The structural checker emits structured diagnostic states rather than a single scalar confidence score. It first checks whether any mapped question entity and any candidate answer node are connected in the evidence subgraph. If a path exists, EVIDENCEFIRST extracts the shortest path as the evidence chain passed to the reader. A path shorter than the question-type minimum is labeled as a short chain rather than accepted as complete. If no adequate path exists, EVIDENCEFIRST records the concrete failure state. A missing entity means that an entity required by the question is absent from the evidence graph. A disconnected pair means that relevant endpoint entities are present but lie in separate components. A bridge gap means that a partial path reaches one side of the answer entity but lacks the relation needed to connect the reasoning chain. A short chain means that a connection exists but is too shallow for the inferred question type. Empty evidence means that KG retrieval returned no usable support and the reader must fall back to passages.

Representative saved 2Wiki records show why graph visibility and answer correctness should be inspected separately. A correct answer can carry a missing-entity graph label, while a structurally connected candidate path can mismatch a posthoc gold-chain proxy. These divergences motivate audit and routing states over the saved graph.

Diagnostic repair is localized to the reported gap. For missing entities, EVIDENCEFIRST extracts triples about the missing entity from the retrieved passages and adds them to the evidence graph. For disconnected pairs, EVIDENCEFIRST processes at most five entity pairs and extracts bridging triples from the same passage context. After each repair stage, the structural checker is run again and the new chain length, gap type, and repair statistics are appended to D . When the graph remains incomplete, the reader still receives the available triples together with local passages, while D records the residual gap. This design avoids presenting a repaired answer as if the graph had been complete from the start. Repair uses only the candidate passage and graph context available to the system; it does not use gold answers or gold support facts at test time. The structural check and saved-state replay are deterministic functions of the saved triples and logs; graph repair itself may use the same temperature-zero LLM extraction backend as KG construction, so replayability is defined over saved artifacts.

3.4 Grounded Reading, Selection, and Canonicalization

The answer reader receives the selected evidence triples and local passages associated with the current evidence route. The reader is instructed to produce a concise answer rather than a long explanation. A final selection step decides whether to use the evidence-grounded postprocessed answer, and a lightweight canonicalizer normalizes title casing, entity aliases, yes/no outputs, and short answer forms when supported by the available context.

This stage is intentionally downstream of structural path-state checking. Reader context can compensate for imperfect graph evidence, but the saved route

remains attached to the output. As a result, EVIDENCEFIRST can distinguish cases where the graph was structurally complete but the reader failed from cases where generation was attempted after repair or fallback. Local passage context is enabled in the main EVIDENCEFIRST setting and disabled only in the reader-context ablation, so the ablation isolates how much answer realization depends on text grounding beyond graph triples.

The selector is conservative. It accepts the refined answer when the original answer is malformed, when the refined answer is normalization-equivalent to the original answer, when the refined answer is a shorter supported form, or when it repairs a question-type mismatch such as a yes/no placeholder for an entity question. Otherwise, the KG-grounded answer is preserved. The diagnostic record therefore includes not only graph states, but also whether the final answer was selected from the refined evidence-grounded path.

4 Experimental Setup

Datasets. We evaluate on two multi-hop QA settings. HotpotQA-1000 is a 1000-question distractor-style evaluation set with repaired contexts from the HotpotQA validation split [25]. 2WikiMultiHopQA-500 is a 500-question sample from 2WikiMultiHopQA [12]; each example contains natural-language passages and Wikidata-derived evidence triples, entity ids, evidence ids, and answer ids. The 2Wiki sample contains 5000 question-context pairs and 3346 unique passage bodies for graph-based baselines. During inference, EVIDENCEFIRST uses the question, candidate passages, and constructed graph fields only; support annotations, evidence ids, answer ids, and gold answers are loaded only after prediction for offline diagnostics.

Both benchmarks combine Web-style text with linked entities and require multi-hop evidence composition; 2Wiki additionally exposes Wikidata-derived support identifiers for offline diagnostics. They instantiate a fixed-context Web QA setting rather than live Web search, so the experiments test evidence-state visibility over available passages and graph evidence, not crawling, freshness, or changing page content. The objective is evidence-state analysis rather than leaderboard optimization, so we evaluate representative subsets that support detailed replay inspection and audit analysis. Each evaluated question produces a full per-example artifact with a constructed KG, state trace, repair provenance, and reader input, and the artifact bundle participates in three offline audits including a 100-case blind packet; the audit protocol, rather than generation throughput, bounds the evaluation scale. Subset-scale evaluation is consistent with recent multi-hop RAG practice: IRCOT evaluates on 500-question samples per dataset [23]. More importantly, every baseline comparison in this paper is paired over identical question IDs on the same subsets with shared rescoring, so subset choice bounds the generalization scope of our conclusions but does not disadvantage any compared system; separate recomputation confirms that the audit-risk trends are stable across 250, 500, and 1000 examples.

Evaluation questions. The evaluation is organized around four Web Knowledge-RAG questions. First, can a pre-reader graph-state contract be attached to answer artifacts while preserving usable answer quality? Second, do the saved states create replayable audit queues without rerunning generation? Third, do graph-state labels provide useful routing causes as structural audit labels? Fourth, how sensitive are the answer results and evidence visibility to reader context and postprocessing under explicit protocol boundaries?

Baselines. On HotpotQA-1000, we compare with Naive RAG, IRCOT [23], A-RAG [4], LightRAG [9], MS GraphRAG [5], and HopRAG strict [16]. On 2WikiMultiHopQA-500, the completed baselines are Naive RAG, IRCOT, A-RAG, LightRAG, and HopRAG strict. Main tables include completed full-sample artifacts with independent strict rescoring. For HotpotQA, paired claims use the strict response-only HopRAG row; the artifact bundle also keeps the legacy HopRAG extraction file (0.5470 EM / 0.6498 F1) for reproducibility.

Protocol and context-budget boundary. The main tables use saved artifacts with explicit context rows: HotpotQA reports the end-to-end run and a selector-sensitivity row, a deterministic re-application of the answer selector over the same saved artifact, while 2Wiki reports a context-rich reader-full audit artifact, where the saved reader input contains the full 2Wiki context rather than only local top- k passages, and a stricter local-context stress row. A separate 2Wiki visibility audit measures how much annotated support is visible in reader inputs or saved graph state after prediction. The answer-quality rows therefore test whether replayable observability telemetry can be attached to usable answer-quality artifacts under recorded, not uniform, evidence conditions.

Metrics and reporting boundary. We use exact match (EM) and token F1 with standard QA normalization: lowercasing, article removal, punctuation removal, and whitespace collapse. HotpotQA reports both the end-to-end saved run and a deterministic selector sensitivity row over saved artifacts; the selector changes 21 predictions and 12 EM labels. Pairwise reporting uses shared question IDs, bootstrap CIs, exact McNemar tests for EM, and explicit labels for exploratory subgroup rows or compound ablation switches. The 2Wiki visibility audit uses support-title recall and gold evidence-entity coverage only after prediction; it does not supply support facts, evidence ids, or gold answers to inference.

Implementation settings. The reported runs use the repository’s OpenAI-style backend with qwen-plus on a DashScope-compatible endpoint unless a run script explicitly overrides it. Answering and repair use temperature 0; KG construction artifacts are cached and loaded for the reported evaluations. The main pipeline uses max-iter 3, local passage top- $k = 5$, and at most five disconnected endpoint pairs per repair pass. The selector and 2Wiki canonicalizer are deterministic. Verifier thresholds and ablation switches are fixed before evaluation. Per-example JSONL artifacts record cache status, ablation flags, path-state status, gap labels, repair/fallback steps, and answer-selection decisions, so audit summaries can be recomputed without new generation.

Artifacts and reproducibility boundary. Every reported result is computed from completed full-sample artifacts with independent strict rescoring. The offline no-LLM verifier checks pinned SHA256, size, row-count, and schema fingerprints; independently rescores prediction/gold rows; and recomputes the main saved summaries without fresh generation. The quick presence check, `python scripts/verify_paper_artifacts.py -check-only`, checks 40 required files, 45 fingerprints, 12 2Wiki source-field files, and 17,600 rescored rows.

5 Results

5.1 Evidence-State Validity

The first result is that EVIDENCEFIRST creates inspectable graph states before answer scoring. On HotpotQA, 56.8% of examples end with a checked candidate path, repair is attempted on 63.7%, and 43.49% of attempted repairs end with a structurally connected recheck; on 2Wiki, the corresponding rates are 47.8%, 60.0%, and 22.0%. Internal label-consistency checks over saved fields show that missing-entity, disconnected, and short-chain labels satisfy their structural invariants on both datasets; for the HotpotQA empty-gap marker used to identify checked paths, 568 of 577 rows (98.44%) also carry `chain_complete=true`, with the remaining nine rows carrying an empty final gap label despite an incomplete chain. Thus, the method records whether an answer came from checked-path, repaired, or residual incomplete evidence before answer scoring.

Table 2 calibrates the states as observability metadata rather than semantic correctness certificates. Structural consistency checks whether saved fields satisfy the rule that produced the state. Gold-proxy precision checks alignment with a corresponding posthoc gold-evidence condition; proxy match reports a 100-example three-proxy offline audit; and blind match reports a separate 100-case blind packet judged without gold answers, EM/F1, baselines, or prior proxy judgments. Low match rates are therefore a calibration signal for structural metadata, not a contradiction of the state trace.

Table 2. 2Wiki structural audit-state analysis. Structural consistency is computed over all 500 saved examples; proxy match is from the 100-row offline proxy plausibility audit; blind match is from a separate 100-case isolated blind audit packet.

Audit state	N	EM/F1	Struct. cons.	Gold-proxy prec.	Proxy match	Blind match
complete	239	0.703 / 0.768	1.000	0.151	0.118	0.677
missing entities	217	0.696 / 0.767	1.000	0.687	0.743	0.571
short chain	33	0.636 / 0.704	1.000	0.455	0.400	0.200
disconnected	11	0.455 / 0.535	1.000	0.091	0.182	0.546

The asymmetry in Table 2 is informative. When the checker reports a deficiency, the label is comparatively trustworthy: missing-entity gold-proxy precision is 0.687 and proxy match is 0.743. When it declares completeness, the label is structurally consistent but semantically optimistic, with gold-proxy precision of 0.151. This asymmetry matches the intended use: deficiency states trigger repair and inspection routing, while the checked-path label certifies only structural connectivity, never semantic sufficiency. Missing-entity EM remains close to

complete EM on both 2Wiki (0.696 vs. 0.703) and HotpotQA (0.560 vs. 0.590), which indicates reader-context compensation rather than state-label failure; the reader-context ablation in Table 8 supports the same interpretation by reducing 2Wiki EM from 0.690 to 0.184 when text grounding is removed.

For HotpotQA, the saved per-state outcomes show the same boundary without a gold-proxy audit: complete-label rows have 0.590 EM / 0.676 F1, missing-entity rows have 0.560 / 0.658, short-chain rows have 0.441 / 0.567, and disconnected rows are too few for interpretation ($N = 3$). These values support the same claim boundary: state labels separate structural routing conditions, not standalone correctness levels.

Because 2Wiki includes support facts and evidence triples, we use them only after prediction to audit evidence visibility and gold-chain alignment. Table 3 shows that the reader-full audit condition keeps all annotated support titles visible, while local top- k retrieval traces and A-RAG read chunks expose substantially less annotated support. The saved EVIDENCEFIRST KG covers 76.57% of gold evidence entities, recalls 47.78% of gold endpoint pairs, and exactly matches 4.20% of gold relations. Thus the checked/repaired states are operational evidence states, while exact gold-chain recovery remains a separate audit target.

Table 3. 2Wiki evidence-visibility audit. Support-title recall and evidence entity coverage use 2Wiki annotations only after prediction; passage rows scan visible text, while the saved-KG row scans graph nodes and edges.

Evidence object	N	Support recall	Entity cov.	EM/F1
EVIDENCEFIRST reader-full input	500	1.0000	0.9468	0.6900/0.7581
Naive BM25 top-5 input	500	0.6020	0.6685	–
A-RAG actual read chunks	500	0.2955	0.3326	–
A-RAG search-found upper bound	500	0.5885	0.6085	–
EVIDENCEFIRST saved KG	500	–	0.7657	0.6900/0.7581

The structural states are selective audit strata. Checked-path cases reach 70.29% EM and 76.81% F1, while disconnected-gap cases fall to 45.45% EM and 53.51% F1 and short-chain cases fall to 63.64% EM and 70.35% F1. Missing-entity cases remain close to complete cases, and the gold-chain audit shows low relation recall. The resulting interpretation is that the states expose structural strata for auditing and repair without replacing semantic evidence validation.

5.2 Replayable Audit Utility

The second result is that saved states turn GraphRAG outputs into recomputable audit queues. The no-LLM verifier replays the artifact boundary without fresh generation, checks source-field hygiene for 2Wiki inference artifacts, and independently rescores saved prediction/gold rows before trusting reported EM/F1 values. More importantly, the same record lets an operator recompute which answers came from checked paths, repaired paths, residual gaps, or fallback states rather than rerunning the model.

Table 4 summarizes the saved audit signals as global rankers. The composite audit-risk score combines answer-selection and graph-state signals; inside a 20%

review budget, saved audit queues have 67.0% error rate on HotpotQA and 40.0% on 2Wiki, corresponding to $1.53\times$ and $1.29\times$ lift over the base error rates. Graph-only signals are modest as global rankers, but they still provide typed causes that answer-only selection cannot supply.

Table 4. Triage utility of saved audit signals. AUC detects answer errors; top-20 error is the error rate among the highest-risk 20% examples, and lift is relative to each dataset’s base error rate.

Signal	Hotpot AUC	Hotpot top-20 err.	Hotpot lift	2Wiki AUC	2Wiki top-20 err.	2Wiki lift
Composite audit risk	0.6321	0.6700	1.5297	0.5787	0.4000	1.2903
Answer not selected	0.6077	0.7000	1.5982	0.5754	0.4000	1.2903
Graph-only audit	0.5434	0.5250	1.1986	0.5230	0.3500	1.1290
Chain incomplete	0.5321	0.5050	1.1530	0.5144	0.2900	0.9355

Table 5 isolates the incremental utility of the evidence-state fields. Answer selection alone is the stronger scalar risk signal, and graph state alone is weak. Combining the two nevertheless improves HotpotQA AUC from 0.6077 to 0.6321 and slightly improves 2Wiki AUC from 0.5754 to 0.5787. The interpretation is therefore complementary rather than competitive: graph state is not a replacement for answer-selection confidence, but it adds routing evidence and typed failure causes beyond that confidence.

Table 5. Incremental utility of evidence-state signals for answer-error triage. Values reuse the saved audit-signal scores from Table 4.

Signal	Hotpot AUC	2Wiki AUC
Answer selection only	0.6077	0.5754
Graph state only	0.5434	0.5230
Selection + graph state	0.6321	0.5787

We note explicitly that at the top-20% operating point, the composite score does not dominate answer selection alone: HotpotQA top-20 error is 0.670 versus 0.700, and 2Wiki is identical at 0.400. The incremental value of graph state is therefore not a stronger scalar ranker. It appears in ranking quality over the full score range (HotpotQA AUC 0.6077 to 0.6321) and, operationally, in the typed failure causes of Section 5.3 that a scalar selection signal cannot provide. We also check that the audit-risk signal is not an artifact of a single sample size: recomputing the composite score on 250-, 500-, and 1000-example prefixes of the same HotpotQA artifact yields AUC 0.632, 0.655, and 0.632 with top-20% error 0.640, 0.690, and 0.670, so the reported audit trends are stable beyond roughly 500 examples.

5.3 Routing Utility and Case Studies

The third result is that graph state is useful as a routing cause, not as a standalone correctness oracle. Inside answer-selection queues, Table 7 shows two positive conditional contrasts: selected unrepaired incomplete Hotpot cases are 26.6 error points worse than selected checked-path cases, and non-selected residual-gap 2Wiki cases are 33.3 error points worse than non-selected cases without residual gaps. These contrasts make the graph state actionable: route checked-path failures

to reader or selector inspection, residual gaps to graph repair, and empty/fallback states to retrieval or context review.

Table 6. Failure-routing interpretation of saved evidence states.

Evidence state	Suggested inspection
Checked failure	Reader / selector
Residual gap	Graph construction / bridge repair
Missing entity	Entity extraction
Fallback	Retrieval / context review

Table 7. Conditional routing utility of graph state inside answer-selection queues. Error-rate deltas use Wilson/Newcombe 95% intervals.

Dataset	Conditional queue contrast	N_1/N_2	Error ₁	Error ₂	Δ error [95% CI]
Hotpot	selected checked path $\rightarrow$ selected unrepaired incomplete	468/51	0.3419	0.6078	0.2660[0.1220, 0.3947]
2Wiki	non-selected no residual gap $\rightarrow$ non-selected residual gap	137/11	0.3942	0.7273	0.3331[0.0285, 0.5249]

The same logic is visible in individual records. Figure 2 expands a repaired-path HotpotQA case into a compact before/after replay view: the Ted Kooshian question is routed from no checked path to bridge repair, then to a rechecked length-2 path and the correct answer EGOT. The point of the case study is the route, not the answer string: the saved fields identify the pre-repair graph state, the repair action, and the post-repair route without regenerating the answer. This example illustrates how evidence-state traces can support post-hoc inspection and routing decisions.




Fig. 2. Replayable Repair Trace for a representative saved case. The question, graph state, repair, reader route, and answer are all recoverable from saved fields, so the before/after repair decision can be audited without regeneration.

Table 8 reports four operational removals on both datasets. Two switches are intentionally compound: disabling structural checking also disables gap-driven repair, and disabling reader context also disables passage-based answer refinement. The largest answer-quality drop comes from removing reader-context and refinement, which reduces HotpotQA EM from 0.5620 to 0.2360 and 2Wiki EM from 0.6900 to 0.1840. This result shows that graph triples alone do not provide enough surface evidence for robust answer realization.

The checker removal should be read as loss of evidence-state observability, not as an isolated EM ablation. When chain checking and repair are disabled, the reader can still answer from unchecked context, so answer-level effects are mixed and HotpotQA even increases slightly. The 0.6-point EM difference between the full protocol and the no-checker variant is not statistically distinguishable (exact

McNemar $p = 0.659$ over 1000 paired examples). The operational loss is clearer: the system no longer records whether the reader saw a checked path, repaired path, residual graph gap, or fallback evidence, and it loses the typed state needed to route graph repair versus reader inspection.

Table 8. Compound operational removals for EVIDENCEFIRST. Full rows use the selector-variant protocol; ablations use explicit switches, with two compound rows labeled in the text rather than interpreted as isolated causal effects.

Variant	HotpotQA-1000		2Wiki-500	
	EM	F1	EM	F1
Full EVIDENCEFIRST selector protocol	0.5620	0.6564	0.6900	0.7581
w/o state checker + repair	0.5680	0.6647	0.6780	0.7579
w/o graph repair	0.5540	0.6580	0.6500	0.7424
w/o reader context + refinement	0.2360	0.2899	0.1840	0.2289
w/o answer refinement	0.5510	0.6485	0.6900	0.7782

5.4 QA Quality Preservation

The final result is an answer-quality preservation boundary. On HotpotQA-1000, the saved EVIDENCEFIRST end-to-end row reaches 55.0% EM / 64.8% F1 and the deterministic selector-sensitivity row reaches 56.2% EM / 65.6% F1, within the operating range of HopRAG strict (54.1% / 64.4%) and A-RAG (51.7% / 62.7%). Paired full-set intervals against HopRAG cross zero (selector ΔEM 0.021, 95% CI $[-0.008, 0.050]$), so these rows are not a new SOTA claim; the systems claim rests on the added state trace, replay, and routing utility. The HotpotQA bridge subset remains exploratory but positive (ΔEM 0.058, 95% CI $[0.016, 0.100]$).

On 2WikiMultiHopQA-500, the context-rich reader-full audit row reaches 69.0% EM / 75.8% F1, near LightRAG (67.4% / 77.2%) and A-RAG (66.6% / 73.2%); the local-context stress row reaches 50.4% / 59.5% and marks the stricter context boundary. Reader-full EVIDENCEFIRST is positive against HopRAG (ΔEM 0.126, 95% CI $[0.084, 0.166]$) but not separated from LightRAG or A-RAG, confirming that context budget must be reported explicitly.

5.5 Retrieval-Noise Probe

Under live Web retrieval (Tavily API, 100 HotpotQA questions, seed=123, top-5 snippets, pipeline unchanged), the state distribution shifts sharply: checked-path rates drop from 29.0% to 20.0%, repaired-path rates drop from 31.0% to 13.0%, and residual-gap rates rise from 39.0% to 67.0%; the single benchmark fallback row (1.0%) becomes checked under live retrieval. Of the 29 benchmark-checked questions, 17 degrade to residual and 3 trigger successful bridge repair; of the 31 benchmark-repaired questions, 17 degrade to residual but 8 remain repaired. The degradation and the survival of the repair route confirm that the state contract correctly responds to retrieval noise. The probe is exploratory (N=100), tests state-contract behavior not answer quality, and is replayable from archived search results.

6 Discussion

The main implication is that GraphRAG systems need observability, not only better retrieval and generation. Databases expose transaction logs, and distributed systems expose traces; fixed-context Web Knowledge-RAG similarly needs an evidence-state trace that records what graph condition the reader saw before an answer was produced. EVIDENCEFIRST supplies such a trace for four operational states: checked path, repaired path, residual gap, and passage fallback.

Why replayability matters. Traditional logs record what happened, but replayable evidence-state traces also allow operators to reconstruct why a routing decision was made and independently verify audit conclusions without regenerating answers. This distinction matters for production GraphRAG because repeated execution may no longer be identical after model, prompt, index, or Web-context changes. By storing the pre-reader state, repair route, fallback tag, and selection decision, EVIDENCEFIRST turns failure analysis from passive log inspection into reproducible post-deployment audit.

The evidence-state trace is most valuable for failure attribution. A wrong answer with a checked path should first be inspected as a reader or selector failure. A wrong answer with a residual gap should first be inspected as graph construction or bridge-repair failure. A fallback case should be inspected as a retrieval or context-visibility failure. This routing story also explains the mixed ablations: removing the checker does not have to reduce EM if the reader can still answer from unchecked text, but it removes the state needed to decide which subsystem should be repaired.

The state labels should not be read as semantic proof that the evidence is sufficient. The 2Wiki gold-chain and blind audits show that structural labels only partially align with human or gold-proxy notions of support. This is a feature of the claim boundary: EVIDENCEFIRST makes evidence state visible and replayable; it does not replace semantic verification, professional adjudication, or uniform-budget Web retrieval.

EVIDENCEFIRST’s contribution is orthogonal to retrieval and generation quality: it neither significantly improves nor degrades the answer-quality operating range and should be evaluated as an observability layer, not as a replacement for the reader or retriever, the same way a tracing system is not judged by whether it speeds up the traced service.

7 Threats to Validity

Protocol validity. The experiments define a fixed-context Web Knowledge-RAG setting with recorded, but not uniform, context budgets. Our experiments support graph-state visibility under fixed-context Web QA settings: every reported row has a saved evidence condition, the 2Wiki table separates reader-full and local-context rows, and the verifier recomputes the main claims without fresh generation. Section 5.5 provides only a small-sample retrieval-noise probe; deployment-scale

evaluation under open-domain retrieval and dynamic Web environments remains future work.

Diagnostic validity. The checker is structural rather than semantic, depends on lexical mapping and coarse answer-type filters, and only partially aligns with gold 2Wiki paths. The gold-chain, proxy, and blind audits calibrate the saved states as routing metadata while keeping semantic evidence validation separate.

Deployment validity. The artifact boundary has practical consequences: 2Wiki lacks a completed MS GraphRAG run, ablations are partly compound, and the selector row is deterministic postprocessing. Passage fallback remains lightly exercised: 5/1000 HotpotQA and 0/500 2Wiki benchmark rows, and the retrieval-noise probe of Section 5.5 also does not trigger it, since Web search rarely returns unusable text for these questions; the fallback route is therefore preserved in the state contract and verifier but not stress-tested. Future work should add professional gap-label adjudication, fallback stress tests, semantic repair-success analysis, and deployment-scale uniform-budget Web retrieval runs.

8 Conclusion

We presented EVIDENCEFIRST, a replayable evidence-state auditing and routing layer for fixed-context Web Knowledge-RAG. By shifting audit logging from only the answer string to the retrieved evidence graph, EVIDENCEFIRST exposes checked paths, repaired paths, residual gaps, passage fallback, and answer-selection state as recomputable audit metadata. The result is not a claim that structural graph state certifies correctness, but a practical evidence-state trace that may help operators replay, inspect, and analyze GraphRAG failures after deployment.

References

1. Buneman, P., Khanna, S., Tan, W.: Why and where: A characterization of data provenance. In: Proceedings of the 8th International Conference on Database Theory (ICDT). pp. 316–330 (2001)
2. Chen, J., Lin, H., Han, X., Sun, L.: Benchmarking large language models in retrieval-augmented generation. In: Proceedings of AAAI. vol. 38, pp. 17754–17762 (2024)
3. Cheney, J., Chiticariu, L., Tan, W.C.: Provenance in databases: Why, how, and where. *Foundations and Trends in Databases* **1**(4), 379–474 (2009)
4. Du, M., Xu, B., Zhu, C., Wang, S., Wang, P., Wang, X., Mao, Z.: A-RAG: Scaling agentic retrieval-augmented generation via hierarchical retrieval interfaces. arXiv preprint arXiv:2602.03442 (2026)
5. Edge, D., et al.: From local to global: A graph RAG approach to query-focused summarization. arXiv preprint arXiv:2404.16130 (2024)
6. Es, S., James, J., Anke, L.E., Schockaert, S.: RAGAS: Automated evaluation of retrieval augmented generation. In: Proceedings of EACL System Demonstrations. pp. 150–158 (2024)

7. Friel, R., Belyi, M., Sanyal, A.: RAGBench: Explainable benchmark for retrieval-augmented generation systems. arXiv preprint arXiv:2407.11005 (2024)
8. Gao, T., Yen, H., Yu, J., Chen, D.: Enabling large language models to generate text with citations. In: Proceedings of EMNLP. pp. 6465–6488 (2023)
9. Guo, Z., Xia, L., Yu, Y., Ao, T., Huang, C.: LightRAG: Simple and fast retrieval-augmented generation. arXiv preprint arXiv:2410.05779 (2024)
10. Gutiérrez, B.J., Shu, Y., Gu, Y., Yasunaga, M., Su, Y.: HippoRAG: Neurobiologically inspired long-term memory for large language models. In: Advances in Neural Information Processing Systems (NeurIPS) (2024)
11. Han, H., Wang, Y., Shomer, H., Guo, K., Ding, J., Lei, Y., Halappanavar, M., Rossi, R.A., Mukherjee, S., Tang, X., et al.: Retrieval-augmented generation with graphs (GraphRAG). arXiv preprint arXiv:2501.00309 (2025)
12. Ho, X., Nguyen, A.D., Sugawara, S., Aizawa, A.: Constructing A multi-hop QA dataset for comprehensive evaluation of reasoning steps. In: Proceedings of COLING. pp. 6609–6625 (2020)
13. Lewis, P., et al.: Retrieval-augmented generation for knowledge-intensive NLP tasks. In: Advances in Neural Information Processing Systems (NeurIPS). pp. 9459–9474 (2020)
14. Li, M., Zou, J., Lv, X., Zhang, C., Zhou, G.: S2G-RAG: Structured sufficiency and gap judging for iterative retrieval-augmented QA. arXiv preprint arXiv:2604.23783 (2026), accepted to ACL 2026 Main Conference
15. Li, Z., et al.: StructRAG: Boosting knowledge intensive reasoning of LLMs via inference-time hybrid information structurization. arXiv preprint arXiv:2410.08815 (2024)
16. Liu, H., Wang, Z., Chen, X., Li, Z., Xiong, F., Yu, Q., Zhang, W.: HopRAG: Multi-hop reasoning for logic-aware retrieval-augmented generation. arXiv preprint arXiv:2502.12442 (2025)
17. Mohan, C., Haderle, D., Lindsay, B., Pirahesh, H., Schwarz, P.: ARIES: A transaction recovery method supporting fine-granularity locking and partial rollbacks using write-ahead logging. ACM Transactions on Database Systems **17**(1), 94–162 (1992)
18. Rashkin, H., et al.: Measuring attribution in natural language generation models. Computational Linguistics **49**(4), 777–840 (2023)
19. Ru, D., et al.: RAGChecker: A fine-grained framework for diagnosing retrieval-augmented generation. arXiv preprint arXiv:2408.08067 (2024)
20. Saad-Falcon, J., Khattab, O., Potts, C., Zaharia, M.: ARES: An automated evaluation framework for retrieval-augmented generation systems. In: Proceedings of NAACL. pp. 338–354 (2024)
21. Sarthi, P., Abdullah, S., Tuli, A., Khanna, S., Goldie, A., Manning, C.D.: RAPTOR: Recursive abstractive processing for tree-organized retrieval. In: International Conference on Learning Representations (ICLR) (2024)
22. Sigelman, B.H., et al.: Dapper, a large-scale distributed systems tracing infrastructure. Tech. rep., Google, Inc. (2010)
23. Trivedi, H., Balasubramanian, N., Khot, T., Sabharwal, A.: Interleaving retrieval with chain-of-thought reasoning for knowledge-intensive multi-step questions. In: Proceedings of ACL. pp. 10014–10037 (2023)
24. Wallat, J., Heuss, M., de Rijke, M., Anand, A.: Correctness is not faithfulness in RAG attributions. arXiv preprint arXiv:2412.18004 (2024)
25. Yang, Z., et al.: HotpotQA: A dataset for diverse, explainable multi-hop question answering. In: Proceedings of EMNLP. pp. 2369–2380 (2018)